

C Programming in Linux

David Haskins

David Haskins

C Programming in Linux

C Programming in Linux

2nd edition

© 2014 David Haskins & bookboon.com

ISBN 978-87-403-0543-2


Contents

About the author, David Haskins	7
Introduction	9
Setting up your System	12
1 Hello World	14
1.1 Hello Program 1	14
1.2 Hello Program 2	15
1.3 Hello Program 3	18
1.4 Hello Program 4	20
1.5 Hello World conclusion	23
2 Data and Memory	24
2.1 Simple data types?	24
2.2 What is a string?	28

Free eBook on Learning & Development

By the Chief Learning Officer of McKinsey

Download Now



Click on the ad to read more

2.3	What can a string “mean”	29
2.4	Parsing a string	32
2.5	Data and Memory – conclusion	34
3	Functions, pointers and structures	36
3.1	Functions	36
3.2	Library Functions	38
3.3	A short library function reference	39
3.4	Data Structures	41
3.5	Functions, pointers and structures – conclusion	44
4	Logic, loops and flow control	45
4.1	Syntax of C Flow of control	45
4.2	Controlling what happens and in which order	46
4.3	Logic, loops and flow conclusion	56
5	Database handling with MySQL	57
5.1	On not reinventing the wheel	57
5.2	MySQL C API	57



www.sylvania.com

We do not reinvent the wheel we reinvent light.

Fascinating lighting offers an infinite spectrum of possibilities: Innovative technologies and new markets provide both opportunities and challenges. An environment in which your expertise is in high demand. Enjoy the supportive working atmosphere within our global group and benefit from international career paths. Implement sustainable ideas in close cooperation with other specialists and contribute to influencing our future. Come and join us in reinventing light every day.

Light is OSRAM

OSRAM
SYLVANIA



Click on the ad to read more

6	Graphics with GD library	61
6.1	Generating binary content	61
6.2	Using TrueType Fonts	63
6.3	GD function reference	65
7	Apache C modules	69
7.1	Safer C web applications	69
7.2	Adding some functionality	72
7.3	Apache Modules Conclusion	73
8	The Ghost project	74
8.1	A PHP web site generator project	74
12	Conclusion	80

**Deloitte.**Discover the truth at www.deloitte.ca/careers

© Deloitte & Touche LLP and affiliated entities.



About the author, David Haskins

I was born in 1950 in Chelsea, London, but grew up in New Zealand returning to England in 1966. I have worked in the computer industry since 1975 after a couple of years as a professional drummer.

My first experience was five years as a mainframe hardware engineer for Sperry Univac (now Unisys) followed by 14 years as an analyst programmer with British Telecom in London.

While engaged in a complex task of converting large quantities of geographical data (map coordinate references) I discovered the joys of C – its speed and efficiency. That was in 1985 and I have been a fan of C ever since.

Since 1994 I have been a senior lecturer at the Faculty of Computing, Information Systems and Mathematics at Kingston University, London. This is a mostly technical university that evolved from a former polytechnic college with a long tradition of aeronautical engineering.

I am engaged mainly in teaching many computer languages and internet systems design to a large and multicultural student body.

Most of my academic research and commercial consultancy has been involved with spatial systems design and the large data volumes and necessary processing efficiency concerns has led me to concentrate on C and C++. My teaching web site is at www.ubiubi.org which shows some of this material.

A keen Open Systems enthusiast, I have exclusively centred all my teaching on the Linux platform since 2002 and Kingston University is well advanced in delivering dual boot facilities for all its student labs.



I am a keen swimmer and in 2009 completed the annual Lorne *Pier-to-Pub* race in Victoria, Australia which is the largest open-sea swimming race in the world where 4,500 people of all ages swim each January as the shark-spotting planes fly overhead.

When not teaching I am a keen vegetable gardener and amateur musician, playing in jazz groups and in Scottish bagpipe bands. I play the drums, the great highland bagpipe, the clarinet, the guitar and the piano.

Introduction

Why learn the C language?

Because the C language is like Latin – it is finite and has not changed for years. C is tight and spare, and in the current economic climate we will need a host of young people who know C to keep existing critical systems running.

C is built right into the core of Linux and Unix. The design idea behind Unix was to write an operating system in C so all you needed to port it to a new architecture was a C compiler. Linux is essentially the success story of a series of earlier attempts to make a PC version of Unix.

A knowledge of C is now and has been for years a pre-requisite for serious software professionals and with the recent popularity and maturity of Open Systems this is even more true. The terseness and perceived difficulty of C saw it being ousted from university teaching during the late 1990s in favour of Java but there is a growing feeling amongst some teaching communities that Java really is not such a good place to start beginners.

Students paradoxically arrive at colleges knowing less about computing than they did ten years ago as programming is seen as too difficult for schools to teach. Meanwhile the body of knowledge expected of a competent IT professional inexorably doubles every few years.

Java is commonly taught as a first language but can cause student confusion as it is in constant flux, is very abstract and powerful, and has become too big with too many different ways to do the same thing. It also is a bit “safe” and insulates students from scary experiences, like driving with air-bags and listening to headphones so you take less care. The core activity of writing procedural code within methods seems impenetrable to those who start from classes and objects.

So where do we start? A sensible place is “at the beginning” and C is as close as most of us will ever need to go unless we are becoming hardware designers. Even for these students to start at C and go further down into the machine is a good idea.

C is like having a very sharp knife which can be dangerous, but if you were learning to be a chef you would need one and probably cut yourself discovering what it can do. Similarly C expects you to know what you are doing, and if you don’t it will not warn before it crashes.

A knowledge of C will give you deep knowledge of what is going on beneath the surface of higher-level languages like Java. The syntax of C pretty-well guarantees you will easily understand other languages that came afterwards like C++, Java, Javascript, and C#.

C gives you access to the heart of the machine and all its resources at a fine-grained bit-level.

C has been described as like “driving a Porsche with no brakes” – and because it is fast as well this can be exhilarating. C is often the only option when speed and efficiency is crucial.

C has been called “dangerous” in that it allows low-level access to the machine but this scariness is exactly what you need to understand as it gives you respect for the higher-level languages you will use.

Many embedded miniaturised systems are all still written in C and the machine-to-machine world of the invisible internet for monitoring and process control often uses C.

Hopefully this list of reasons will start you thinking that it might be a good reason to have a go at this course.

References

The C Programming Language – Second Edition – Kernighan and Richie
ISBN 0-13-11-362-8

The GNU C Library Free Software Foundation C Manual
<http://www.gnu.org/software/libc/manual/>

MySQL C library
<http://dev.mysql.com/doc/refman/5.5/en/index.html>

The GD C library for graphics
<http://www.libgd.org/Documentation>

APXS – the APache eXtenSion tool
<http://httpd.apache.org/docs/2.2/programs/apxs.html>

Apache
<http://httpd.apache.org/docs/2.2/developer/>

“The Apache Modules Book” Nick Kew, Prentice Hall
ISBN 0-13-240967-4

A **Source Code Zip File Bundle** is supplied with this course which contains all the material described and a Makefile.

The teaching approach

I began university teaching later in life after a career programming in the telecommunications industry.

My concern has been to convey the sheer fun and creativity involved in getting computers to do what you want them to do and always try to give useful, practical, working examples of the kinds of things students commonly tell me they want to do.

Learning a language can be a dry, boring affair unless results are immediate and visible so I tend to use the internet as the input-output channel right from the start.

I prefer teaching an approach to programming which is deliberately “simple” using old-fashioned command-line tools and editors and stable, relatively unchanging components that are already built-in to Unix and Linux distributions such as Suse, Ubuntu and Red Hat.

This is in response to the growing complexity of modern Integrated Development Environments (IDEs) such as Developer Studio, Netbeans and Eclipse which give students an illusion that they know what they are doing but generate obfuscation.

My aim is to get students confident and up to speed quickly without all the nightmare associated with configuring complex tool chains. It is also essentially a license-free approach and runs on anything.

With this fundamental understanding about what is really going on you can progress on to use and actually understand whatever tools you need in your career.

In order to give a sense of doing something real and useful and up to date, the focus is on developing visible and effectively professional-quality web-server and client projects to put on-line, using:

Apache Web server and development libraries.

C language CGI programs (C programming using the “make” utility).

C language Apache modules.

MySQL server with C client library interfaces.GD graphics library with C interfaces.

Incidental use of CSS, (X)HTML, XML, JavaScript, Ajax.

This course has been designed for and lab-tested by first and second year Computer Science Students at Kingston University, London UK.

Setting up your System

This book presumes you are using the Linux operating system with either a KDE or Gnome desktop. Specific set-up instructions for common architectures are at <http://www.ubiubi.org/CProgrammingInLinux/>

If you are using the KDE desktop you will have Konqueror or Dolphin as the File Manager and kate or kedit for an editor

In Gnome you would probably use Nautilus and gedit

You need to be familiar with the idea of doing some things as “super user” so that you have access permission to copy or edit certain files. This is normally done by prefacing the Linux command with “sudo” and providing the password, as in this example:

```
“sudo cp hello3 /srv/www/cgi-bin/hello3”
```

which copies the file “hello3” to the area where the Apache server locates common gateway interface or cgi programs.

SIMPLY CLEVER

ŠKODA

We will turn your CV into an opportunity of a lifetime

Do you like cars? Would you like to be a part of a successful brand? We will appreciate and reward both your enthusiasm and talent. Send us your CV. You will be surprised where it can take you.

Send us your CV on www.employerforlife.com

Click on the ad to read more

In KDE “kdesu konqueror” would open a file manager as super user.

In Gnome “gnomesu nautilus” would open a file manager as super user.

You will need to have installed the following packages:

package	Ubuntu	Open Suse
C development libraries	build-essential	Base Development (pattern)
Apache web server	apache2	Web and LAMP Server (pattern)
Apache development libraries	apache2-prefork-dev	apache2-devel
MySQL server, client and development libraries	mysql-server libmysqlclient-dev	libmysqlclient-devel
GD and development libraries	libgd2-xpm libgd2-xpm-dev	gd gd-devel

Throughout the text you will see references to the folder **cgi-bin**. The location of this will vary between Linux distributions. By default this folder used for web programs is:

OpenSuse: **/srv/www/cgi-bin**
 Ubuntu: **/usr/lib/cgi-bin**


To place programs there you need superuser rights, so it may be better to create a folder inside your own home/***/public_html/cgi-bin directory and change the **ScriptAlias** and associated **Directory** references inside the Apache configuration files (OpenSuse) **/etc/apache2/default-server.conf** or (Ubuntu) **/etc/apache2/sites-available/default**.

1 Hello World

1.1 Hello Program 1

Using the File Manager (in KDE, Konqueror or in Gnome, Nautilus) create a new directory somewhere in your home directory called something appropriate for all the examples in this book, perhaps “Programming_In_Linux” without any spaces in the name.

Open an editor (in KDE, kate, or in Gnome, gedit) and type in (or copy from the supplied source code zip bundle) the following:



Save the text as **chapter1_1.c** in the new folder you created in your home directory.


Open a terminal window and type: **gcc -o hello chapter1_1.c**
to compile the program into a form that can be executed.

Now type “ls -l” to list the details of all the files in this directory. You should see that chapter1_2.c is there and a file called “hello” which is the compiled C program you have just written.

Now type: **./hello**
to execute, or run the program and it should return the text:

“Hello you are learning C!!”

If this worked, congratulations, you are now a programmer!




1.2 Hello Program 2

Taking this example a stage further, examine the start of the program at the declaration of the entry point function: `int main(int argc, char *argv[])`

In plain English this means:

The function called “main”, which returns an integer, takes two arguments, an integer called “argc” which is a count of the number of command arguments then *argv[] which is a list or array of pointers to strings which are the actual arguments typed in when you run the program from the command line.




Let's rewrite the program to see what all this means before we start to panic.

The advertisement features a woman teacher smiling and interacting with two young students (a boy and a girl) who are looking at a laptop screen. The background is yellow with orange swirling graphics. The e-Learning for Kids logo is in the top left. In the bottom right, there's a green oval containing text about the organization's impact.

About e-Learning for Kids Established in 2004, e-Learning for Kids is a global nonprofit foundation dedicated to fun and free learning on the Internet for children ages 5 - 12 with courses in math, science, language arts, computers, health and environmental skills. Since 2005, more than 15 million children in over 190 countries have benefitted from eLessons provided by EFK! An all-volunteer staff consists of education and e-learning experts and business professionals from around the world committed to making difference. eLearning for Kids is actively seeking funding, volunteers, sponsors and courseware developers; get involved! For more information, please visit www.e-learningforkids.org.

The number 1 MOOC for Primary Education
• Free Digital Learning for Children 5-12
• 15 Million Children Reached



Save the text as **chapter1_2.c** in the same folder.

Open a terminal window and type:

gcc -o hello2 chapter1_2.c to compile the program into a form that can be executed.

Now type **ls -l** to list the details of all the files in this directory. You should see that **chapter1_2.c** is there and a file called **hello2** which is the compiled C program you have just written.

Now type **./hello2** to execute, or run the program and it should return the text:

*Hello, you are still learning C!!
Number of arguments to the main function:1
argument number 0 is ./hello2*


We can see that the name of the program itself is counted as a command line argument and that the counting of things in the list or array of arguments starts at zero not at one.

Now type **./hello2 my name is David** to execute the program and it should return the text:

*Hello, you are still learning C!!
Number of arguments to the main function:5
argument number 0 is ./hello2*

*argument number 1 is my
argument number 2 is name
argument number 3 is is
argument number 4 is David*

So, what is happening here? It seems we are reading back each of the character strings (words) that were typed in to run the program.



1.3 Hello Program 3

Lets get real and run this in a web page. Make the extra change adding the first output printf statement “Content-type:text/plain\n\n” which tells our server what kind of MIME type is going to be transmitted.

Compile using **gcc -o hello3 chapter1_3.c** and copy the compiled file hello3 to your public_html/cgi-bin directory (or on your own machine as superuser copy the program to /srv/www/cgi-bin (OpenSuse) or /usr/lib/cgi-bin (Ubuntu)).



Open a web browser and type in the URL <http://localhost/cgi-bin/hello3?david+haskins> and you should see that web content can be generated by a C program.


Cynthia | AXA Graduate

**AXA Global
Graduate Program**

Find out more and apply

redefining / standards **AXA**

A green oval at the bottom right contains a hand cursor icon pointing to the text "Click on the ad to read more".




1.4 Hello Program 4

A seldom documented feature of the function signature for “main” is that it can take **three** arguments and the last one we will now look at is `char *env[]` which is also a list of pointers to strings, but in this case these are the **system environment variables** available to the program at the time it is run



Compile with **gcc -o hello4 chapter1_4.c** and as superuser copy the program to **/srv/www/cgi-bin** (OpenSuse) or **/usr/lib/cgi-bin** (Ubuntu). You can run this from the terminal where you compiled it with **./hello4** and you will see a long list of environment variables. In the browser when you enter <http://localhost/cgi-bin/hello4> you will see a different set altogether.



The screenshot shows a Mozilla Firefox window displaying the output of a CGI script. The title bar reads "File Edit View History Bookmarks Tools Help". The address bar shows "http://localhost/cgi-bin/hello4?david+has". The main content area displays the following text:

```
Hello, you are still learning C!!  
Number of arguments to the main function:3  
argument number 0 is /srv/www/cgi-bin/hello4  
argument number 1 is david  
argument number 2 is haskins  
Environment variables:  
env[0] = HTTP_HOST=localhost  
env[1] = HTTP_USER_AGENT=Mozilla/5.0 (X11; U; Linux i686; en-GB; rv:1.9.0.3) Gecko/20080  
env[2] = HTTP_ACCEPT=text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8  
env[3] = HTTP_ACCEPT_LANGUAGE=en-gb,en;q=0.5  
env[4] = HTTP_ACCEPT_ENCODING=gzip,deflate  
env[5] = HTTP_ACCEPT_CHARSET=ISO-8859-1,utf-8;q=0.7,*;q=0.7  
env[6] = HTTP_KEEP_ALIVE=300  
env[7] = HTTP_CONNECTION=keep-alive  
env[8] = PATH=/bin:/sbin:/usr/bin:/usr/sbin  
env[9] = SERVER_SIGNATURE=<address>Apache/2.2.8 (Linux/SUSE) Server at localhost Port 80  
  
env[10] = SERVER_SOFTWARE=Apache/2.2.8 (Linux/SUSE)  
env[11] = SERVER_NAME=localhost  
env[12] = SERVER_ADDR=127.0.0.1  
env[13] = SERVER_PORT=80  
env[14] = REMOTE_ADDR=127.0.0.1  
env[15] = DOCUMENT_ROOT=/srv/www/htdocs  
env[16] = SERVER_ADMIN=[no address given]  
env[17] = SCRIPT_FILENAME=/srv/www/cgi-bin/hello4  
env[18] = REMOTE_PORT=19911  
env[19] = GATEWAY_INTERFACE=CGI/1.1  
env[20] = SERVER_PROTOCOL=HTTP/1.1  
env[21] = REQUEST_METHOD=GET  
env[22] = QUERY_STRING=david+haskins  
env[23] = REQUEST_URI=/cgi-bin/hello4?david+haskins  
env[24] = SCRIPT_NAME=/cgi-bin/hello4
```

The status bar at the bottom of the browser window says "Done".

We will soon find out that **QUERY_STRING** is an important environment variable for us in communicating with our program and in this case we see it has a value of “david+haskins” or everything after the “?” in the URL we typed. It is a valid way to send information to a **common gateway interface** (CGI) program like hello4 but we should restrict this to just one string. In our case we have used a “+” to join up two strings. If we typed: “david haskins” the browser would translate this so we would see:

```
QUERY_STRING=david%20haskins
```

We will learn later how complex sets of input values can be transmitted to our programs.

I joined MITAS because
I wanted **real responsibility**

The Graduate Programme
for Engineers and Geoscientists
www.discovermitas.com



Month 16

I was a construction supervisor in the North Sea advising and helping foremen solve problems

Real work
International opportunities
Three work placements




Click on the ad to read more

1.5 Hello World conclusion

We have seen that a simple program with a tiny bit of input and some output is in fact extremely powerful in that it reveals and exposes the inner workings of a great deal of our computer.

Even though we have just begun we have encountered many of the key concepts we will use over and over again:

- functions and arguments
- Numbers (integers) and character strings as data types
- Lists or arrays
- Loops using “for” and “while”

We have made a deliberate big leap from writing a program that runs simply in a “terminal screen” to one which will be visible over the internet in a browser.

The reason for this is that the process of writing programs that interact with users in windowing systems like Windows, Gnome or KDE is extremely complex and not something you will be asked very often to do .

The internet browser has become the de facto interface mode for almost everything we do these days so we might as well understand using it from the start.

In all the successive chapters we will follow this model: starting off with some basic technique then applying it to a web-based system.

In practice there is not much real-world C common gateway interface programming going on but there is a great deal of C and C++ based code running as Apache modules and Microsoft IIS ISAPI Dlls. Perhaps not many know that much of Ebay is written in C / C++.

Why? It is as fast as things get and their business with the bargain snipers in a global real-time market needs this lightning fast core, so there is no other way to get that performance.

2 Data and Memory

2.1 Simple data types?

When we write programs we have to make decisions or assertions about the nature of the world as we declare and describe variables to represent the kinds of things we want to include in our information processing.

This process is deeply philosophical; we make **ontological** assertions that this or that thing exists and we make **epistemological** assertions when we select particular data types or collections of data types to use to describe the attributes of these things. Heavy stuff with a great responsibility and not to be lightly undertaken.

As a practical example we might declare something that looks like the beginnings of a database record for geography.



Here we are doing the following:

- asserting that all the character strings we will ever encounter in this application will be 255 limited to characters so we **define** this with a **preprocessor** statement – these start with #.
- assert that towns are associated with counties, and counties are associated with countries some hierarchical manner.
- assert that the population is counted in whole numbers – no half-people.
- assert the location is to be recorded in a particular variant (WGS84) of the convention of describing spots on the surface of the world in latitude and longitude that uses a decimal fraction for degrees, minutes, and seconds.

Each of these statements allocates **memory** within the **scope** of the function in which it is declared. Each **data declaration** will occupy an amount of memory in **bytes** and give that bit of memory a label which is the **variable name**. Each data type has a specified size and the **sizeof()** library function will return this as an integer. In this case 3×256 characters, one integer, and two floats. The exact size is machine dependent but probably it is 780 bytes.

Outside the function in which the data has been declared this data is inaccessible – this is the **scope** of declaration. If we had declared outside the `main()` function it would be **global** in scope and other functions could access it. C lets you do this kind of dangerous stuff if you want to, so be careful.

The advertisement features a background image of modern skyscrapers. On the left, the IE business school logo is displayed. The central text highlights that 93% of MIM students are working in their sector 3 months following graduation. Below this, the program title "MASTER IN MANAGEMENT" is prominently displayed. To the left of the main text, there is a list of benefits: "STUDY IN THE CENTER OF MADRID AND TAKE ADVANTAGE OF THE UNIQUE OPPORTUNITIES THAT THE CAPITAL OF SPAIN OFFERS", "PROPEL YOUR EDUCATION BY EARNING A DOUBLE DEGREE THAT BEST SUITS YOUR PROFESSIONAL GOALS", and "STUDY A SEMESTER ABROAD AND BECOME A GLOBAL CITIZEN WITH THE BEYOND BORDERS EXPERIENCE". To the right, detailed information about the program is provided: Length: 10 MONTHS, Av. Experience: 1 YEAR, Language: ENGLISH / SPANISH, Format: FULL-TIME, Intakes: SEPT / FEB. At the bottom, three boxes highlight "5 SPECIALIZATIONS PERSONALIZE YOUR PROGRAM", "#10 WORLDWIDE MASTER IN MANAGEMENT FINANCIAL TIMES", and "55 NATIONALITIES IN CLASS". Contact information includes the website www.ie.edu/master-management, email mim.admissions@ie.edu, and social media links for Facebook, Twitter, and LinkedIn. The footer also includes the text "Follow us on IE MIM Experience".



Click on the ad to read more

Generally we keep a close eye on the scope of data, and pass either **read-only copies**, or **labelled memory addresses** to our data to parts of the programs that might need to do work on it and even change it. These labelled memory addresses are called **pointers**.

We are using for output the **printf** family of library functions (**sprintf** for creating strings, **fprintf** for writing to files etc.) which all use a common **format string** argument to specify how the data is to be represented.

- %c character
- %s string
- %d integer
- %f floating point number etc.


The remaining series of variables in the arguments are placed in sequence into the format string as specified.

In C it is a good idea to **initialise** any data you declare as the contents of the memory allocated for them is not cleared but may contain any old rubbish.

Compile with: **gcc -o data1 chapter2_1.c -lc**

Output of the program when called with : **./data1**

```
Town name: Guildford population:66773
County: Surrey
Country: Great Britain
Location latitude: 51.238598 longitude: -0.566257
char = 1 byte int = 4 bytes float = 4 bytes
memory used:780 bytes
```



2.2 What is a string?

Some programming languages like Java and C++ have a **string data type** that hides some of the complexity underneath what might seem a simple thing.


An essential attribute of a character string is that it is a series of individual character elements of indeterminate length.

Most of the individual characters we can type into a keyboard are represented by simple numerical ASCII codes and the C data type **char** is used to store character data.

Strings are stored as **arrays** of characters ending with a NULL so an array must be large enough to hold the sequence of characters plus one. Remember array members are always counted from zero.

In this example we can see 5 individual characters declared and initialised with values, and an empty character array set to “”.

Take care to notice the difference between single quote marks ‘ used around characters and double quote marks “ used around character strings.





Compile with: **gcc -o data2 chapter2_2.c -lc**

Output of the program when called with : **./data2**

david

2.3 What can a string "mean"

Anything at all – **name** given to a **variable** and its **meaning** or its **use** is entirely in the mind of the beholder. Try this




Compile with: **gcc -o data3 chapter2_3.c -lc**


As superuser copy the program to your public_html/cgi-bin directory (or /srv/www/cgi-bin (OpenSuse) or /usr/lib/cgi-bin (Ubuntu)).

In the browser enter: <http://localhost/cgi-bin/data3?red>

what you should see is this:



Or if send a parameter of anything at all you will get surprising results:



What we are doing here is **using** the string parameter `argv[1]` as a background colour code inside an **HTML** body tag. We change the Content-type specification to `text/html` and miraculously now our program is generating HTML content. A language being expressed inside another language. Web browsers understand a limited set of colour terms and colours can be also defined hexadecimal codes such as `#FFFFFF` (white) `#FF0000` (red) `#00FF00` (green) `#0000FF` (blue).

This fun exercise is not just a lightweight trick, the idea that one program can generate another in another language is very powerful and behind the whole power of the internet. When we generate HTML (or XML or anything else) from a **common gateway interface program** like this we are creating **dynamic content** that can be linked to live, changing data rather than **static** pre-edited web pages. In practice most web sites have a mix of dynamic and static content, but here we see just how this is done at a very simple level.

Throughout this book we will use the browser as the preferred interface to our programs hence we will be generating HTML and binary image stream web content purely as a means to make immediate the power of our programs. Writing code that you peer at in a terminal screen is not too impressive, and writing window-type applications is not nearly so straightforward.

In the past 5 years we have drilled around
95,000 km
—that's more than **twice** around the world.

Who are we?
We are the world's leading provider of reservoir characterization, drilling, production, and processing technologies to the oil and gas industry.

Who are we looking for?
We offer countless opportunities in the following domains:

- Operations
- Research, Engineering, and Manufacturing
- Geoscience and Petrotechnical
- Commercial and Business

We're looking for high-energy, self-motivated graduates with vision and integrity to join our team.

careers.slb.com

What will you be?

Schlumberger

In practice most of the software you may be asked to write will be running on the web so we might as well start with this idea straight away. Most web applications involve **multiple languages** too such as CSS, (X)HTML, XML, JavaScript, PHP, JAVA, JSP, ASP, .NET, SQL. If this sounds frightening, don't panic. A knowledge of C will show you that many of these languages, which all perform different functions, have a basis of C in their syntax.

2.4 Parsing a string

The work involved in extracting meaning or valuable information from some kind of input string is called “parsing”. We will now build another fun internet-callable CGI program to demonstrate the power in our hands.




Compile with: **gcc -o data4 chapter2_4.c -lc**

As superuser copy the program to /srv/www/cgi-bin (OpenSuse) or /usr/lib/cgi-bin (Ubuntu).

In the browser enter:

<http://localhost/cgi-bin/data4?red:blue:5:5>

what you should see is this:



In this program we take argv[1] which here is **yellow:blue:5:5**: and **parse** it using the library function **strtok** which chops the string into tokens separated by an arbitrary character ':' and use these tokens as strings to specify colours and integer numbers to specify the row and cell counts of a table.

The function **atoi** converts an string representation of a integer to an integer ("1" to 1).

The function **strtok** is a little odd in that the first time you call it with the string name you want to parse, then on subsequent calls the first parameter is changed to NULL.

The **for(...)** loop mechanism was used to do something a set number of times.

The HTML terms introduced were:

<html> <body> <table> <tr> table row <td> table data cell

2.5 Data and Memory – conclusion

We have used some simple data types to represent some information and transmit input to a program and to organise and display some visual output.

We have used HTML embedded in output strings to make output visible in a web browser.

As an exercise try this:



Write a program to put into your **/public_html/cgi-bin** folder which can be called in a browser with the **name** of a **sports team** or a **country** and a series of **colours** specified perhaps as hexadecimals e.g. ff0000 = red (rrggb) used for the team colours or map colours, and which displays something sensible.

My version looks like this:

The advertisement features a photograph of a young woman in a black dress shaking hands with a man in a suit in a formal setting with portraits on the wall. The text includes:
Excellent Economics and Business programmes at:
university of
groningen
AACSB ACCREDITED
“The perfect start
of a successful,
international career.”
CLICK HERE
to discover why both socially
and academically the University
of Groningen is one of the best
places for a student to be
www.rug.nl/feb/education



Click on the ad to read more



3 Functions, pointers and structures

3.1 Functions

The entry point into all our programs is called main() and this is a **function**, or a **piece of code** that does something, usually returning some value. We structure programs into functions to stop them become long unreadable blocks of code than cannot be seen in one screen or page and also to ensure that we do not have repeated identical chunks of code all over the place. We can call **library functions** like printf or strtok which are part of the C language and we can call our own or other peoples functions and libraries of functions. We have to ensure that the appropriate header file exists and can be read by the preprocessor and that the source code or compiled library exists too and is accessible.

As we learned before, the **scope** of data is restricted to the **function** in which it was declared, so we use **pointers** to data and blocks of data to pass to functions that we wish to do some work on our data. We have seen already that strings are handled as pointers to arrays of single characters terminated with a NULL character.




In this example we can repeatedly call the function “doit” that takes two integer arguments and returns the result of some mathematical calculation.

Compile: gcc -o func1 chapter3_1.c -lm

Copy to cgi-bin: cp func1 /home/david/public_html/cgi-bin/func1

(You should be using the Makefile supplied or be maintaining a Makefile as you progress, adding targets to compile examples as you go.)



American online LIGS University

is currently enrolling in the
Interactive Online **BBA, MBA, MSc,**
DBA and PhD programs:

- ▶ enroll **by September 30th, 2014** and
- ▶ **save up to 16%** on the tuition!
- ▶ pay in 10 installments / 2 years
- ▶ Interactive Online education
- ▶ visit www.ligsuniversity.com to find out more!

Note: LIGS University is not accredited by any nationally recognized accrediting agency listed by the US Secretary of Education.
[More info here.](#)

Click on the ad to read more

The result in a browser looks like this called with “func1?5:5”:

In this case the arguments to our function are sent as **copies** and are not modified in the function but used.

If we want to actual modify a variable we would have to send its pointer to a function.



We send the address of the variable ‘result’ with **&result**, and in the function do it we **de-reference** the pointer with ***result** to get at the float and change its value, outside its scope inside main. This gives identical output to chapter3_1.c.

3.2 Library Functions

C contains a number of built-in functions for doing commonly used tasks. So far we have used **atoi**, **printf**, **sizeof**, **strtok**, and **sqrt**. To get full details of any built-in library function all we have to do is type for example:

man 3 atoi

and we will see all this:

```

david@linux-bxso:~ - Shell - Konsole
Session Edit View Bookmarks Settings Help
Atoi(3)                               Linux Programmer's Manual      Atoi(3)
NAME
    atoi, atol, atoll, atoq - convert a string to an integer
SYNOPSIS
#include <stdlib.h>

int atoi(const char *nptra);
long atol(const char *nptra);
long long atoll(const char *nptra);
long long atoq(const char *nptra);

Feature Test Macro Requirements for glibc (see feature_test_macros(7)):
atoll(): _BSD_SOURCE || _SVID_SOURCE || _XOPEN_SOURCE >= 600 ||
_ISOC99_SOURCE; or cc -std=c99

DESCRIPTION
The atoi() function converts the initial portion of the string pointed
to by nptra to int. The behavior is the same as
    strtol(nptra, (char **) NULL, 10);
except that atoi() does not detect errors.

The atol() and atoll() functions behave the same as atoi(), except that
they convert the initial portion of the string to their return type of
long or long long. atoq() is an obsolete name for atoll().

RETURN VALUE
The converted value.

CONFORMING TO
SVr4, POSIX.1-2001, 4.3BSD, C99. C89 and POSIX.1-1996 include the
functions atoi() and atol() only. atoq() is a GNU extension.

NOTES
The non-standard atoq() function is not present in libc 4.6.27 or glibc
2, but is present in libc5 and libc 4.7 (though only as an inline func-
tion in <stdlib.h> until libc 5.4.44). The atoll() function is present
in glibc 2 since version 2.0.2, but not in libc4 or libc5.

SEE ALSO
atof(3), strtod(3), strtol(3), strtoul(3)

Manual page atoi(3) line 1


```

Which pretty-well tells you everything you need to know about this function and how to use it and variants of it. Most importantly it tells you which **header file** to include.

3.3 A short library function reference

Full details of all the functions available can be found at:

<http://www.gnu.org/software/libc/manual/>



There is no point in learning about library functions until you find you need to do something which then leads you to look for a function or a library of functions that has been written for this purpose. You will need to understand the function signature – or what the argument list means and how to use it and what will be returned by the function or done to variables passed as pointers to functions.

3.4 Data Structures

Sometimes we wish to manage a set of variable as a group, perhaps taking all the values from a database record and passing the whole record around our program to process it. To do this we can group data into structures.

This program uses a struct to define a set of properties for something called a player. The main function contains a declaration and instantiation of an array of 5 players. We pass a pointer to each array member in turn to a function to rank each one. This uses a switch statement to examine the first letter of each player name to make an arbitrary ranking. Then we pass a pointer to each array member in turn to a function that prints out the details.



DON'T EAT YELLOW SNOW

What will your advice be?

Some advice just states the obvious. But to give the kind of advice that's going to make a real difference to your clients you've got to listen critically, dig beneath the surface, challenge assumptions and be credible and confident enough to make suggestions right from day one. At Grant Thornton you've got to be ready to kick start a career right at the heart of business.

Sound like you? Here's our advice: visit GrantThornton.ca/careers/students

Scan here to learn more about a career with Grant Thornton.


 **Grant Thornton**
An instinct for growth™

© Grant Thornton LLP. A Canadian Member of Grant Thornton International Ltd





The results are shown here, as usual in a browser:




This is a very powerful technique that is quite advanced but you will need to be aware of it. The idea of **structures** leads directly to the idea of **classes** and **objects**.

We can see that using a struct greatly simplifies the business task of passing the data elements around the program to have different work done. If we make a change to the definition of the struct it will still work and we simply have to add code to handle new properties rather than having to change the argument lists or signatures of the functions doing the work.

The definition of the structure does not actually create any data, but just sets out the formal shape of what we can instantiate. In the main function we can express this instantiation in the form shown creating a list of sequences of data elements that conform to the definition we have made.

You can probably see that a struct with additional functions or methods is essentially what a class is in Java, and this is also the case in C++. **Object Oriented languages** start here and in fact many early systems described as “object oriented” were in fact just built using C language structs.

If you take a look for example, at the Apache server development header files you will see a lot of structs for example in this fragment of httpd.h:



Dont worry about what this all means – just notice that this is a very common and very powerful technique, and the design of data structures, just like the design of database tables to which it is closely related are the core, key, vital task for you to understand as a programmer.

You make the philosophical decisions that the **world is like this** and can be modelled in this way. A heavy responsibility – in philosophy this work is called **ontology** (what exists?) and **epistemology** (how we can know about it?). I bet you never thought that this was what you were doing!

3.5 Functions, pointers and structures – conclusion

We have used some simple data types to represent some information and transmit input to a program and to organise and display some visual output.

We have used HTML embedded in output strings to make output visible in a web browser.

We have learned about creating libraries of functions for reuse.

We have learning about data structures and the use of pointers to pass them around a program.

Exercise:

Using C library functions create a program that:

- opens a file in write mode,
- writes a command line argument to the file
- closes the file
- opens the file in read mode
- reads the contents
- closes the file
- prints this to the screen

This will give you experience with finding things out, looking for suitable library functions, and finding examples on-line or from a book.

4 Logic, loops and flow control

4.1 Syntax of C Flow of control

We can use the following C constructs to control program execution.

When we can count our way through a sequence or series:

```
for( initial value; keep on until ; incremental change )
    { do this; and this; and this; }
```

When we are waiting for some condition to change:

```
while( this is true )
    { do this; and this; and this; }
```

or if we want to do something at least once then test:

```
do { do this; and this; and this; }
while( this is true )
```

..... Alcatel-Lucent 

www.alcatel-lucent.com/careers

What if you could build your future and create the future?

One generation's transformation is the next's status quo. In the near future, people may soon think it's strange that devices ever had to be "plugged in." To obtain that status, there needs to be "The Shift".



Click on the ad to read more

When we have a single option to test:

```
if( this is true )
    { do this; and this; and this; }
else
    { do this; and this; and this; }
```

When we have more options to test:

```
if( this is true )
    { do this; and this; and this; }
else if ( this is true )
    { do this; and this; and this; }
else
    { do this; and this; and this; }
```

When we have more options to test based on an integer or single character value:

```
switch( on an integer or character value )
{
    case 0: do this; and this; and this; break;
    case n: do this; and this; and this; break;
    default:do this; and this; and this; break;
}
```

4.2 Controlling what happens and in which order

This part is all about **if**, and **then**, and **else** and **true** and **false** – the nuts and bolts of how we express and control the execution of a program. This can be very dry and dusty material so to make it more understandable we are going to solve a problem you are going to need to solve to do any interactive web work of any complexity.


We will build something we can use in order to provide something like the functionality that can be obtained from typical **getParameter("ITEM1")** method in Java servlets or **\$_REQUEST["ITEM1"]** function in PHP.

In Chapter 1 we saw that environment variables can be accessed by the implicit argument to the main function. We can also use the library function **getenv()** to request the value of any named environment variable.




Here we display the **QUERY_STRING** which is what the program gets as the entire contents of an HTML form which contains NAME=VALUE pairs for all the named form elements.


An HTML form by default uses the **GET** method which transmits all form data back to the program or page that contains the form unless otherwise specified in an action attribute. This data is contained in the **QUERY_STRING** as a series of variable = value pairs separated by the & character.



Calling this program in a browser we see a form and can enter some data in the boxes:



And after submitting the form we see:



To make much sense of the QUERY_STRING and find a particular value in it, we are going to have to **parse** it, to chop it up into its constituent pieces and for this we will need some **conditional logic** (if, else etc.) and some **loop** to count through the characters in the variable. A basic function to do this would ideally be created as this is a task you might need to do again and again so it makes sense to have a chunk of code that can be called over again.

In the next example we add this function and the noticeable difference in the output is that we can insert the extracted values into the HTML boxes after we have parsed them. We seem to have successfully created something like a java getParameter() function – or have we?



Maastricht University *Leading in Learning!*

**Join the best at
the Maastricht University
School of Business and
Economics!**

Top master's programmes

- 33rd place Financial Times worldwide ranking: MSc International Business
- 1st place: MSc International Business
- 1st place: MSc Financial Economics
- 2nd place: MSc Management of Learning
- 2nd place: MSc Economics
- 2nd place: MSc Econometrics and Operations Research
- 2nd place: MSc Global Supply Chain Management and Change

Sources: Keuzegids Master ranking 2013; Elsevier 'Beste Studies' ranking 2012; Financial Times Global Masters in Management ranking 2012


**Visit us and find out why we are the best!
Master's Open Day: 22 February 2014**

Maastricht
University is
the best specialist
university in the
Netherlands
(Elsevier)

www.mastersopenday.nl




Click on the ad to read more




Have a good long look at chapter4_2.c and try it out with characters other than A-Z a-z or numerals and you will see something is not quite right. There is some kind of encoding going on here!

If I were to type **DAVID !!!** into the first field:



I get this result:



A **space** character has become a **+** and **!** has become **%21**.

This encoding occurs because certain characters are explicitly used in the transmission protocol itself. The **&** for example is used to separate portions of the QUERY_STRING and the space cannot be sent at all as it is.

Any program wishing to use information from the HTML form must be able to decode all this stuff which will now attempt to do.

The program chapter4_2.c accomplishes what we see so far. It has a main function and a decode_value function all in the same file.

The decode_value function takes three arguments:

- the name of the value we are looking for "ITEM1=" or "ITEM2=".
- the address of the variable into which we are going to put the value if found
- the maximum number of characters to copy

The function looks for the start and end positions in the QUERY_STRING of the value and then copies the characters found one by one to the value variable, adding a NULL character to terminate the string.

**Empowering People.
Improving Business.**

BI Norwegian Business School is one of Europe's largest business schools welcoming more than 20,000 students. Our programmes provide a stimulating and multi-cultural learning environment with an international outlook ultimately providing students with professional skills to meet the increasing needs of businesses.

BI offers four different two-year, full-time Master of Science (MSc) programmes that are taught entirely in English and have been designed to provide professional skills to meet the increasing need of businesses. The MSc programmes provide a stimulating and multi-cultural learning environment to give you the best platform to launch into your career.

- MSc in Business
- MSc in Financial Economics
- MSc in Strategic Marketing Management
- MSc in Leadership and Organisational Psychology


www.bi.edu/master

BI NORWEGIAN BUSINESS SCHOOL

EFMD
EQUIS
ACCREDITED




Click on the ad to read more



It looks like we are going to have to do some serious work on this decode_value package so as this is work we can expect to do over and over again it makes sense to write a **function** that can be **reused**.

First off we can put this function into a separate file called **decode_value.c** and create a file for all the functions we may write called **c_in_linux.h** and compile all this into a **library**. In the Make file we can add:



This looks horrible and complex but all it means is this:

typing “**make all**” will:

```
compile all the *.c files listed in the list OBJ_SRC and into object files *.o  
compile all the object files into a library archive called lib_c_in_linux.a  
compile 4-4 using this new archive.
```

This is the model we will use to keep our files as small as possible and the share-ability of code at its maximum.

We can now have a simpler “main” function file, and files for stuff we might want to write as call-able functions from anywhere really which we do not yet know about. All this is organised into a **library file (*.a for archive)** – these can also be compiled as dynamically loadable **shared objects *.so** which are much like Windows DLLs. This exactly how all Linux software is written and delivered.

For example the MySQL C Application Programmers Interface (API) comprises:

```
all the header files in /usr/include/mysql  
the library file /usr/lib/mysql/libmysqlclient.a
```

What we are doing really is how all of Linux is put together – we are simply adding to it in the same way.

Our **main** file now looks like this:



This code calls the function `decode_value` in the same way but because the library, `c_in_linux.a` was linked in when it was compiled and as it has access to the header file `c_in_linux.h` that lists all the functions in the library it all works properly.


Try to describe the process in **pseudocode** of decoding this QUERY STRING:

```
get the QUERY_STRING  
find the search string "ITEM1=" inside it  
look for the end of the value of "ITEM1="  
copy the value to our "value" variable, translating funny codes such as:  
%21 is ! %23 is #
```

These special codes are generated by the browser so that whatever you put in an HTML form will get safely transmitted and not mess about with the HTTP protocol. There are lots of them and the task for this chapter is to **finish this task off** so that EVERY key on your keyboard works as you think it should!!

Program chapter4_3.c calls this unfinished function decode_value which this far can only cope with the **space** character and ! – it uses **if** and **else** and **for** and the library function **getenv**, **strcpy**, **strlen**, **strstr** in a piece of **conditional logic** in which a string is analysed to find a specific item and this thing then copied into a piece of memory called **value** which has been passed to it.

The result shows the decoded value pasted into the first field;





4.3 Logic, loops and flow conclusion

The most important part of controlling the flow of your program is to have a clear idea about what it is you are trying to do. We have also learned to break our code up into manageable lumps, and started to build and use a library and header file of our own.

Being able to express a process in normal words or **pseudocode** is useful and helps you to break the code into steps.

Use **for loops** to explicitly count through things you know have an ending point.

Use **while** and **do...while** loops to do things until some condition changes.

Use **switch** statements to when integers or single characters determine what happens next.

Use **if** and **else if** and **else** when mutually exclusive things can be tested in a sequence.

Complex sets of **if** and **else** and not (!) conditionals can end up unreadable.

Use braces ({ }) to break it all up into chunks.

Exercise:

A useful **task** now would be to **complete the function decode_value** so you have a useful tool to grab web content from HTML forms decoding all the non alpha-numeric keys on your keyboard.

You will use this exercise again and again so it is worth getting it right.

Need help with your dissertation?

Get in-depth feedback & advice from experts in your topic area. Find out what you can do to improve the quality of your dissertation!

Get Help Now



Go to www.helpmyassignment.co.uk for more info

Helpmyassignment



Click on the ad to read more

5 Database handling with MySQL

5.1 On not reinventing the wheel

It is pretty sensible to not start from scratch at every project so we build on work done by others who came this way. All of the C libraries we are using were written by someone from even lower level bits and pieces and we can access and modify this code should we want to. This is what Open Systems is all about. In practice we simply want to use reliable services, and the first two we are going to use are ubiquitous – databases access and graphical image generation.

5.2 MySQL C API

To access the set of MySQL functions we make sure our compiler can access the header and libraries of MySQL and simply call some mysql functions we have not had to write ourselves.

In this next example we join the database code onto the last work we did and insert the decoded data into a database, and read out what we have stored there so far.

In a new terminal window we type:


```
mysql test
```

then:

```
create table CIL (ITEM1 varchar(255),ITEM2 varchar(255));
```

This creates a table for the program to write to in the test database that should already be created in your MySQL setup.

It looks like this if it all works out OK.



The data types declared **MYSQL**, **MYSQL_RES**, and **MYSQL_ROW** are **preprocessor definitions** that stand for more complex C declarations in mysql.h and all we need to know is how to call them.

The documentation at MySQL provides all the information you will need to do more complex operations.



Understood all that?

In barely 10 lines (highlighted) of C we have inserted our data into a database table and can read it out again, which is pretty painless. With **mysql_init** we obtain a pointer to a data structure of type MYSQL, we can then set this to connect to the test database with **mysql_options** and **mysql_real_connect**, then we execute SQL statements just as we would in a terminal session. The results of a query can be retrieved as a sequence of MYSQL_ROW arrays of strings with **mysql_use_results**. We free up the memory used with **mysql_free_result** and close the database with **mysql_close**.

As is usual with C libraries, you need to be able understand the usually sparse documentation to understand the function calls, and for MySQL 5.1 we can find all this information at:

<http://dev.mysql.com/doc/refman/5.1/en/index.html>

Brain power

By 2020, wind could provide one-tenth of our planet's electricity needs. Already today, SKF's innovative know-how is crucial to running a large proportion of the world's wind turbines.

Up to 25 % of the generating costs relate to maintenance. These can be reduced dramatically thanks to our systems for on-line condition monitoring and automatic lubrication. We help make it more economical to create cleaner, cheaper energy out of thin air.

By sharing our experience, expertise, and creativity, industries can boost performance beyond expectations.

Therefore we need the best employees who can meet this challenge!

The Power of Knowledge Engineering

Plug into The Power of Knowledge Engineering.
Visit us at www.skf.com/knowledge

SKF

The MySQL 5.1 Reference Manual / Connectors and APIs / MySQL C API includes:

mysql_affected_rows()	mysql_insert_id()
mysql_autocommit()	mysql_kill()
mysql_change_user()	mysql_library_end()
mysql_character_set_name()	mysql_library_init()mysql_list_dbs()
mysql_close()	mysql_list_fields()
mysql_commit()	mysql_list_processes()
mysql_connect()	mysql_list_tables()
mysql_create_db()	mysql_more_results()
mysql_data_seek()	mysql_next_result()
mysql_debug()	mysql_num_fields()
mysql_drop_db()	mysql_num_rows()
mysql_dump_debug_info()	mysql_options()
mysql_eof()	mysql_ping()
mysql_errno()	mysql_query()
mysql_error()	mysql_real_connect()
mysql_escape_string()	mysql_real_escape_string()
mysql_fetch_field()	mysql_real_query()
mysql_fetch_field_direct()	mysql_refresh()
mysql_fetch_fields()	mysql_reload()
mysql_fetch_lengths()	mysql_rollback()
mysql_fetch_row()	mysql_row_seek()
mysql_field_count()	mysql_row_tell()
mysql_field_seek()	mysql_select_db()
mysql_field_tell()	mysql_set_character_set()
mysql_free_result()	mysql_set_local_infile_default()
mysql_get_character_set_info()	mysql_set_local_infile_handler()
mysql_get_client_info()	mysql_set_server_option()
mysql_get_client_version()	mysql_shutdown()
mysql_get_host_info()	mysql_sqlstate()
mysql_get_proto_info()	mysql_ssl_set()
mysql_get_server_info()	mysql_stat()
mysql_get_server_version()	mysql_store_result()
mysql_get_ssl_cipher()	mysql_thread_id()
mysql_hex_string()	mysql_use_result()
mysql_info()	mysql_warning_count()
mysql_init()	

6 Graphics with GD library


The ability to generate interesting images dynamically to suit the situation is an enjoyable, challenging, and rewarding type of programming. We will be using Thomas Boutell's GD C library which has been around for many years as an open systems project.

6.1 Generating binary content

Here we create an image and print into it the value of the TEXT= part of the QUERY STRING and set our content type to image/gif. The **gdImageGif** function writes the binary image out to **stdout** which is the output stream instead of to a file, so binary image data is sent back to the browser.



The program produces this output when called with `cgi-bin/gdgraph1?TEXT=DavidHaskins`



While this is not the most fancy or attractive exercise but does at least demonstrate the key principles involved in generating graphical output. With the GD library you can load existing images and generate them in many formats such as **GIF**, **JPEG**, **PNG**, **WMBP** and even create **animated GIFS**.

If you have used PHP development tools you will recognise the GD library functions as they are pretty well identical showing that PHP is a wrapper in this case around the same C library.

The text displayed here is using one of the internal fonts, Tiny, Small, Medium Bold, Large and Giant which are adequate for simple labelling but you can also use TrueType fonts for more attractive output.

Geometric drawing with lines or certain styles, filled and open polygons, and circles, arcs can be created. The main thing to remember is that the origin of an image is the **TOP left hand corner** which might seem unintuitive to anyone who has studied mathematics – quite why this is I have never discovered but can only imagine that the first programmer to do anything in this area happened to not know about graphs in which we think of the origin $x=0, y=0$ as being at the bottom left hand.

Colours are specified as RGB values in the range 0 to 255 so that white is 255,255,255 and red is 255,0,0. To work out fine-grained hues use a graphics tool like GNU GIMP which has colour pickers so you can find out subtle RGB values. Palettes of colours from one image can be used in another and the closest colour in a palette requested or created if there is space for it to be allocated.

6.2 Using TrueType Fonts


In this example we modify the previous program to generate a label using TrueType font.



We call this program as before with `cgi-bin/gdgraph1?TEXT=C+PROGRAMMING+IN+LINUX`

to get this kind of output which you will probably see is more likely to be a useful kind of tool. The location and contents of your systems fonts will vary but the code gives an example:

- OpenSuse /usr/share/fonts/truetype/*.ttf
- Ubuntu. /usr/share/fonts/truetype/ttf-dejavu/*.ttf



To get any good at using a library like GD you have to be prepared to experiment and take a lot of time to understand the function parameters, looking in great detail at the available documentation at:

<http://www.boutell.com/gd>

TURN TO THE EXPERTS FOR SUBSCRIPTION CONSULTANCY

Subscrybe is one of the leading companies in Europe when it comes to innovation and business development within subscription businesses.

We innovate new subscription business models or improve existing ones. We do business reviews of existing subscription businesses and we develop acquisition and retention strategies.

Learn more at [linkedin.com/company/subscrybe/](https://www.linkedin.com/company/subscrybe/) or contact Managing Director Morten Suhr Hansen at mta@subscrybe.dk

SUBSCRYE - to the future



Click on the ad to read more

6.3 GD function reference

A full detailed set of documentation is maintained at: <http://www.boutell.com/gd>

GD contains a wealth of functionality for all kinds of drawing and many formats, as well as TrueType fonts and animated Gif images. A categorised list of functions follows:

Image creation, destruction, loading and saving:

```
gdImageCreate(int sx, int sy)
gdImageCreateFromJpeg(FILE *in)
gdImageCreateFromPng(FILE *in)
gdImageCreateFromGif(FILE *in)
gdImageCreateFromGd(FILE *in)
gdImageCreateFromWBMP(FILE *in)
gdImageDestroy(gdImagePt rim)
void gdImageJpeg(gdImagePt rim, FILE*out, int quality)
void gdImageGif(gdImagePt rim, FILE*out)
void gdImagePng(gdImagePtr im, FILE*out)
void gdImageWBMP(gdImagePtr im, int fg, FILE*out)
void gdImageGd(gdImagePtr im, FILE*out)
```

Drawing Functions:

```
void gdImageSetPixel(gdImagePtr im, int x, int y, int color)
void gdImageLine(gdImagePtr im, int x1, int y1, int x2, int y2, int color)
void gdImageDashedLine(gdImagePtr im, int x1, int y1, int x2, int y2, int color)
void gdImagePolygon(gdImagePtr im, gdPoint Ptr point s, int point sTotal, int color)
void gdImageOpenPolygon(gdImagePtr im, gdPoint Ptr point s, int point sTotal, int color)
void gdImageRectangle(gdImagePtr im, int x1, int y1, int x2, int y2, int color)
void gdImageFilledPolygon(gdImagePtr im, gdPoint Ptr point s, int point sTotal, int color)
void gdImageFilledRectangle(gdImagePtr im, int x1, int y1, int x2, int y2, int color)
void gdImageArc(gdImagePtr im, int cx, int cy, int w, int h, int s, int e, int color)
void gdImageFilledArc(gdImagePtr im, int cx, int cy, int w, int h, int s, int e, int color, int style)
void gdImageFilledEllipse(gdImagePtr im, int cx, int cy, int w, int h, int color)
void gdImageFillToBorder(gdImagePtr im, int x, int y, int border, int color)
void gdImageFill(gdImagePtr im, int x, int y, int color)
void gdImageSetAntiAliased(gdImagePtr im, int c)
void gdImageSetAntiAliasedDontBlend(gdImagePtr im, int c)
void gdImageSetBrush(gdImagePtr im, gdImagePtr brush)
```

```

void gdImageSetTile(gdImagePtr im, gdImagePtr tile)
void gdImageSetStyle(gdImagePtr im, int *style, int styleLength)
void gdImageSetThickness(gdImagePtr im, int thickness)
void gdImageAlphaBlending(gdImagePtr im, int blending)
void gdImageSaveAlpha(gdImagePtr im, int saveFlag)
void gdImageSetClip(gdImagePtr im, int x1, int y1, int x2, int y2)
void gdImageGetClip(gdImagePtr im, int *x1P, int *y1P, int *x2P, int *y2P)

```

Query Functions:

```

int gdImageAlpha(gdImagePtr im, int color)(MACRO)
int gdImageGetPixel(gdImagePtr im, int x, int y)
int gdImageBoundsSafe(gdImagePtr im, int x, int y)
int gdImageGreen(gdImagePtr im, int color)(MACRO)
int gdImageRed(gdImagePtr im, int color)(MACRO)
int gdImageSX(gdImagePtr im)(MACRO)
int gdImageSY(gdImagePtr im)(MACRO)
int gdImageTrueColor(im)(MACRO)

```

Text-handling functions:

```

gdFontPtr gdFontGetSmall(void )
gdFontPtr gdFontGetLarge(void )
gdFontPtr gdFontGetMediumBold(void )
gdFontPtr gdFontGetGiant(void )
gdFontPtr gdFontGetTiny(void )
void gdImageChar(gdImagePtr im, gdFontPtr font, int x, int y, int c, int color)
void gdImageCharUp(gdImagePtr im, gdFontPtr font, int x, int y, int c, int color)
void gdImageString(gdImagePtr im, gdFontPtr font, int x, int y, unsigned char*s, int color)
void gdImageString16(gdImagePtr im, gdFontPtr font, int x, int y, unsigned short *s, int color)
void gdImageStringUp(gdImagePtr im, gdFontPtr font, int x, int y, unsigned char*s, int color)
void gdImageStringUp16(gdImagePtr im, gdFontPtr font, int x, int y, unsigned short*s, int color)

char *gdImageStringFT(gdImagePtr im, int *brect, int fg, char *fontname, double
                      ptsize, double angle, int x, int y, char*string)
char *gdImageStringFTEx(gdImagePtr im, int *brect, int fg, char *fontname, double
                       ptsize, double angle, int x, int y, gdFTString ExtraPtr strex)
char *gdImageStringFTCircle(gdImagePtr im, int cx, int cy, double radius,
                           double textRadius, double fillPortion, char*font, double point s,
                           char*top, char*bottom, int fgcolor)
char *gdImageStringTTF(gdImagePtr im, int *brect, int fg, char *fontname,
                      double ptsize, double angle, int x, int y, char *string)

```

Color-handling functions:

```

int gdImageColorAllocate(gdImagePtr im, int r, int g, int b)
int gdImageColorAllocateAlpha(gdImagePtr im, int r, int g, int b, int a)
int gdImageColorClosest(gdImagePtr im, int r, int g, int b)
int gdImageColorClosestAlpha(gdImagePtr im, int r, int g, int b, int a)
int gdImageColorClosestHWB(gdImagePtr im, int r, int g, int b)
int gdImageColorExact(gdImagePtr im, int r, int g, int b)
int gdImageColorResolve(gdImagePtr im, int r, int g, int b)
int gdImageColorResolveAlpha(gdImagePtr im, int r, int g, int b, int a)
int gdImageColorsTotal(gdImagePtr im)(MACRO)
int gdImageRed(gdImagePtr im, int c)(MACRO)
int gdImageGreen(gdImagePtr im, int c)(MACRO)
int gdImageBlue(gdImagePtr im, int c)(MACRO)
int gdImageGetInterlaced(gdImagePtr im)(MACRO)
int gdImageGetTransparent(gdImagePtr im)(MACRO)
void gdImageColorDeallocate(gdImagePtr im, int color)
void gdImageColorTransparent(gdImagePtr im, int color)
void gdTrueColor(int red, int green, int blue)(MACRO)
void gdTrueColorAlpha(int red, int green, int blue, int alpha)(MACRO)

```



What do you want to do?

No matter what you want out of your future career, an employer with a broad range of operations in a load of countries will always be the ticket. Working within the Volvo Group means more than 100,000 friends and colleagues in more than 185 countries all over the world. We offer graduates great career opportunities – check out the Career section at our web site www.volvologroup.com. We look forward to getting to know you!

VOLVO
AB Volvo (publ)
www.volvologroup.com

VOLVO TRUCKS | RENAULT TRUCKS | MACK TRUCKS | VOLVO BUSES | VOLVO CONSTRUCTION EQUIPMENT | VOLVO PENTA | VOLVO AERO | VOLVO IT
VOLVO FINANCIAL SERVICES | VOLVO 3P | VOLVO POWERTRAIN | VOLVO PARTS | VOLVO TECHNOLOGY | VOLVO LOGISTICS | BUSINESS AREA ASIA



Click on the ad to read more

Resizing functions:

```

void gdImageCopy(gdImagePtr dst, gdImagePtr src, int dstX, int dstY, int srcX, int
srcY, int w, int h)
void gdImageCopyResized(gdImagePtr dst, gdImagePtr src, int dstX, int dstY, int
srcX, int srcY, int destW, int destH, int srcW, int srcH)
void gdImageCopyResampled(gdImagePtr dst, gdImagePtr src, int dstX, int dstY, int
srcX, Int srcY, int destW, int destH, int srcW, int srcH)
void gdImageCopyRotated(gdImagePtr dst, gdImagePtr src, doubledstX, doubledstY,
int srcX, int srcY, int srcW, int srcH, int angle)
void gdImageCopyMerge(gdImagePtr dst, gdImagePtr src, int dstX, int dstY, int srcX,
int srcY, int w, int h, int pct)

void gdImageCopyMergeGray(gdImagePtr dst, gdImagePtr src, int dstX, int dstY, int srcX,
int srcY, int w, int h, int pct)
void gdImagePaletteCopy(gdImagePtr dst, gdImagePtr src)
void gdImageSquareToCircle(gdImagePtr im, int radius)
void gdImageSharpen(gdImagePtr im, int pct)

```

Miscellaneous Functions:

```

int gdImageCompare(gdImagePtr im1, gdImagePtr im2)
gdImageInterlace(gdImagePtr im, int int erlace)
gdFree(void *ptr)

```

In order to use a library like this you will need familiarity with the arguments which are often data types defined within the library itself such as gdimagePtr which is a pointer to some kind of structure containing all the data for an image to be processed or stored. These may all seem unusual but after a while you will begin to get used to the syntax and on-line documentation and begin to see patterns in the complexity.


7 Apache C modules

7.1 Safer C web applications

In real life few web administrators would dream of letting anyone run C programs as CGI content generators because of the risk of crashes and core dumps. However the Apache server is itself written in C and there are simple utilities that come with its development tools that permit you to create code stubs into which you can place your C programs and run them as Apache modules when they are loaded as part of the server and managed safely in a kind of “sand-box”. Here we will take an earlier example and turn it into an Apache module.

A utility called apxs2 is included in the Apache2 development libraries which can be invoked to generate a code stub for a program which can be compiled into a module that is loaded and managed by the Apache web server. These modules can be used to perform a huge variety of tasks but in our case we will do something which is akin the an ISAPI DLL found in the IIS server. The exact location of the apxs2 utility will change according to the Linux distribution you are using but with OpenSuse it runs like this.

In a terminal type: **apxs2 -n labelmaker -g**



```
david@linux-bxso:~ - Shell No. 2 - Konsole
Session Edit View Bookmarks Settings Help
david@linux-bxso:~> /usr/sbin/apxs2 -n labelmaker -g
Creating [DIR] labelmaker
Creating [FILE] labelmaker/Makefile
Creating [FILE] labelmaker/modules.mk
Creating [FILE] labelmaker/mod_labelmaker.c
Creating [FILE] labelmaker/.deps
david@linux-bxso:~>
```

This creates a folder of the name you give it (labelmaker) and a **Makefile**, a **modules.mk** file which can be used by the Make utility, and a file called **mod_labelmaker.c**.

The C file generated is kind of like a Hello World for Apache. It may look like a complex thing but it does supply a long explanatory comment header which is worth reading. The idea is that when Apache starts any modules in a specified location which are configured as needing to be loaded in the server configuration files, will be loaded. The `*_register_hooks` function lists the names and signatures of functions that can be called at specific stages in the Apache server process. In this case if the name `http://localhost/labelmaker` is called this module will be asked to handle whatever happens in the `*_handler` function.

The configuration of the server can be a bit fiddly but in OpenSuse we have to add this to the file

/etc/apache2/sites-available/default


```
<Location /labelmaker>  
    SetHandler labelmaker  
</Location>
```

and in **/etc/config.sys/apache2** we **add** the name of our module labelmaker to long comma-separated list in the line starting

APACHE_MODULES=".....,labelmaker"

Now go to the folder labelmaker and type:

```
sudo apxs2 -c -i mod_labelmaker.c  
sudo /etc/init.d/apache2 restart
```





Call this in a browser like this:



7.2 Adding some functionality

Now we can plug in the work we did for the graphics library in Chapter 6 as a replacement handler function (in the code Chapter7_1.c there are BOTH handlers, one commented out). Note the (highlighted) call to a modified decode_value function that uses the r->args pointer to get the QUERY_STRING rather than getenv(). Also Apache handles the output a bit differently too – get get a pointer to the array of bytes in the image by calling gdImageGifPtr then the ap_rwrite function outputs the data. We have to free the pointer with gdFree after the output call.





7.3 Apache Modules Conclusion

Whilst tricky to write and debug, this is probably the most rewarding and esoteric area where you can do real, commercially useful and safely deployable web content generation. It is easy to see how this example could be extended with parameters for colours and fonts to make a useful web content tool.

There is very little clear simple material about apache modules but start with the on-line documentation at <http://httpd.apache.org/docs/2.2/developer/>

One recent book worth looking at is “The Apache Modules Book” Nick Kew, Prentice Hall.

This e-book
is made with
SetaPDF

SETASIGN

PDF components for PHP developers

www.setasign.com

A green oval button with a white border and a hand cursor icon pointing to it. The text inside the button reads: "Click on the ad to read more".

8 The Ghost project

8.1 A PHP web site generator project

The ability to write short programs in C to automate tedious tasks or to do things that would otherwise take hours of fiddling about with cumbersome tools such as doing mail-merge, is one on the things you will be most pleased you have learned how to do. This project is such a time-saver. Ghost is a lightweight PHP generator for you to customise.


If you find yourself having to build PHP web sites all the time, a quick way to generate all the parameter-passing, decoding, forms building and database management code in one step would be useful. Tools like Ruby on Rails offer such functionality but are infinitely more complex to set up and run and you end up with needing to learn yet another language to go any further.

Probably the best way to start with this tool is to compile and run it. Unzip the ghost.zip source into your public_html folder which creates a folder called ghost. The Makefile contains a target **g1** that compiles and links ghost. So go to **public_html/ghost** and type: **make g1**.

To run the site generator type:

- **./ghost testwebsite data1 data2 data1 data3 data4 data6 data6**
- This will create:
- a folder **public_html/testwebsite**
- a mysql database table called **testwebsite** with text fields data1 data2 data1 data3 data4 data6 data6
- a **testwebsite.css** file
- empty **header.html** and **footer.html** pages
- **index.php** that demonstrates a form handling entry, edit & update, and delete to the database table for the data items specified.

In a browser what you see is this at <http://localhost/~yourname/testwebsite>



The idea behind this is that all the mechanical bits to create and manage the form content are done and can be customised. This screen shot shows the result of submitting one record. The top row is for entering new data, the lower row(s) allow editing or deleting of records. It is a framework that allows you to take and use parts in your own website design.

Let us examine this code in sections.

The first section declares the required data and creates the folder and CSS file.




Next the header.html and footer.html files are generated. These files are loaded by the PHP file and could be used as a generic common header and footers. The CSS file is referenced from the header.html file.



Free eBook on Learning & Development


By the Chief Learning Officer of McKinsey

Download Now



 Click on the ad to read more

Next we create the data base.




The complicated part starts now, of generating a php script. The best way to understand this is to examine the actual output of the program when we view the source of the page in the browser.

The top row is a form with a text box for each column defined in the table generated by running the ghost program.



For each row in the table we now generate a form allowing editing of the data and an anchor link to do a delete operation.



Close examination of the file index.php will allow you to see where all this happens, and to work backward to find where in the ghost.c source code this PHP code is generated. A good idea is to use a highlighter pen on a printout as we are embedding a language (HTML) inside another language (PHP) which is in turn inside another language so very very careful use is made of the escape characters \' to express quotation marks both single and double where necessary to make it all work. This may seem complex – but the speedy prototyping that ghost permits makes it worthwhile to spend time customising the C code so the PHP that you want and the database you want come out the way you want it.

www.sylvania.com

We do not reinvent
the wheel we reinvent
light.

Fascinating lighting offers an infinite spectrum of possibilities: Innovative technologies and new markets provide both opportunities and challenges. An environment in which your expertise is in high demand. Enjoy the supportive working atmosphere within our global group and benefit from international career paths. Implement sustainable ideas in close cooperation with other specialists and contribute to influencing our future. Come and join us in reinventing light every day.


Light is OSRAM

OSRAM
SYLVANIA



Click on the ad to read more

Here is the part of the PHP file index.php which generates the edit or delete rows. The static HTML is highlighted and the other parts are inserted by MySQL PHP function calls.



As you can see a great deal of tedious and repetitive work has been automated. You can move on by modifying the PHP code or go deeper to customise the C program which generates all of it.

I personally use ghost frequently to save time on site-building and this is why I wrote it. I got bored making mistakes writing virtually identical code to decode HTML forms and populate or update databases.

12 Conclusion

Now you have worked through these simple examples I hope you can see why a knowledge of the C language is always going to be a useful and continually practical skill.

After 25 years I still regularly write C programs to do everyday tasks quickly and effectively and once written they form part of a set of durable tools that suit me and which are portable.

This short book contains a number of “tricks” that I have learned over the years for which there is little explicit documentation, and it frankly presupposes a familiarity with Linux. After testing the material with students of a wide range of experience I am confident that an attentive careful student will get all this material working and can start from there to discover the joy of C. Many of my students say “This is much more fun than Java, I can see what is really going on!” which is most gratifying and makes me confident you will find the material useful.

Those interested might ask the publisher to commission me to write a companion volume :

C++ Standard Template Library Programming in Linux

Good luck and happy programming.

David

March 31st 2014